

EECE-5644

INTRO TO MACHINE LEARNING AND PATTERN
RECOGNITION

ASSIGNMENT 3

SAI PRASASTH KOUNDINYA GANDRAKOTA
NUID: 002772719

QUESTION 1

PROCESS

DATA GENERATION

Given that there are 4 classes, and the input X is 3-dimensional, we generate training sets of sizes 100, 200, 500, 1000, 2000 and 5000 along with a test set of size 10000 using the below mean and covariance matrices (varied through trial and error based on keeping the error probability between 10-20%) and the Gaussian Mixture Model with uniform priors:

Handwritten mathematical expressions for mean and covariance matrices and uniform priors:

$$m_1 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} ; m_2 = \begin{bmatrix} 0 \\ 0 \\ 5 \end{bmatrix} ; m_3 = \begin{bmatrix} 0 \\ 5 \\ 0 \end{bmatrix} ; m_4 = \begin{bmatrix} 5 \\ 0 \\ 0 \end{bmatrix}$$
$$C_1 = \begin{bmatrix} 1 & 0 & -2 \\ 0 & 1 & 0 \\ -2 & 0 & 12 \end{bmatrix} ; C_2 = \begin{bmatrix} 6 & 0 & 0 \\ 0 & 0.3 & 0 \\ 0 & 0 & 0.3 \end{bmatrix} ; C_3 = \begin{bmatrix} 1 & 0 & -2 \\ 0 & 0.3 & 0 \\ -2 & 0 & 12 \end{bmatrix} ; C_4 = \begin{bmatrix} 6 & 0 & 2 \\ 0 & 1 & 0 \\ 2 & 0 & 12 \end{bmatrix}$$

Uniform Priors $\rightarrow [0.25, 0.25, 0.25, 0.25]$

NEURAL NETWORK ARCHITECTURE

A 2-layer MLP, with one layer of P hidden perceptrons with the ELU activation function and an output layer with the softmax function applied, is created. Then using 10-fold cross validation, where we split the data into 10 folds, use each fold as a validation data and the rest of the folds apart from the validation data as training data, and then find the best number of perceptrons for the hidden layer (between 1 and 10) based on the minimum classification error probability.

TRAINING THE NEURAL NETWORK

Having identified the optimum number of perceptrons for each training set, we then train 6 different MLPs each with their corresponding data set according to the minimum cross-entropy loss estimation and identify the best weights and biases over 10 reinitializations.

TESTING THE NEURAL NETWORK

Using the respective weights and biases for each MLP we then predict the labels of the test data set and estimate the probability of error along with the confusion matrices, approximating the MAP classification rule.

THEORETICALLY OPTIMUM CLASSIFIER

We train an ideal Bayesian Classifier using the below rule:

$$\text{Error} = \text{Loss Matrix} * \text{Class Posteriors}$$
$$E(Y=C|X) = LM * P(L=C|X)$$

where $X \rightarrow$ input
 $Y \rightarrow$ prediction
 $L \rightarrow$ actual
 $C \rightarrow$ classes (0,1,2,3)
 $E \rightarrow$ error probability
 $LM \rightarrow$ loss matrix

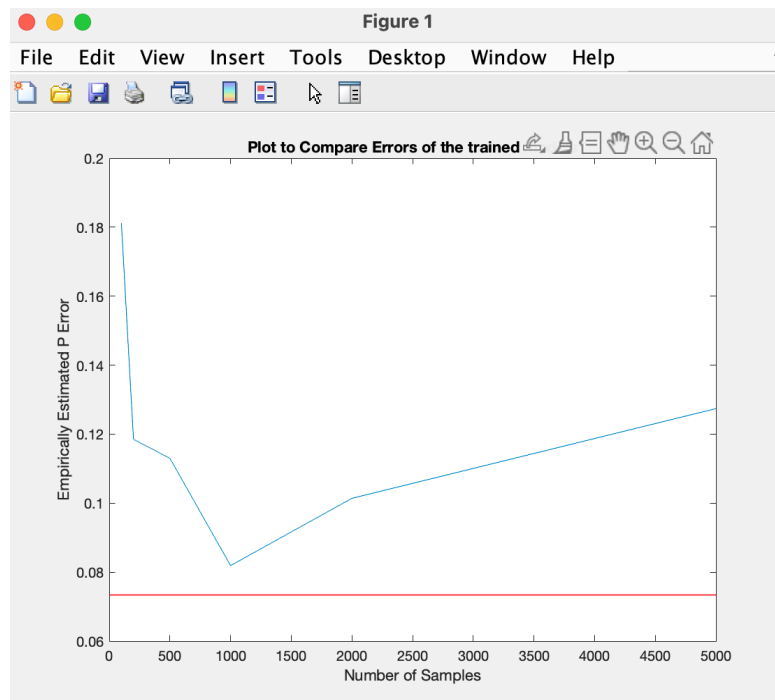
The prediction Y is obtained by minimizing the risk of choosing each class.

RESULTS

Each MLP model along with the optimum number of perceptrons in the hidden layer and their corresponding minimum classification error, final probability error on the test data set and confusion matrix are given below. We see that as the number of perceptrons increases, the error decreases and saturates at a point where we can select the optimum number.

MLP MODEL 1 (100 training samples)			MLP MODEL 4 (1000 training samples)				
-----			-----				
Optimum number of perceptrons = 8			Optimum number of perceptrons = 10				
Min Classification Error = 0.160000			Min Classification Error = 0.211000				
Probability of Error = 0.181200			Probability of Error = 0.081900				
Confusion Matrix =			Confusion Matrix =				
4807	474	9	638	5569	256	1	102
77	776	0	37	102	756	0	32
277	1	565	11	1	0	853	0
163	124	1	2040	257	68	0	2003
MLP MODEL 2 (200 training samples)			MLP MODEL 5 (2000 training samples)				
-----			-----				
Optimum number of perceptrons = 9			Optimum number of perceptrons = 8				
Min Classification Error = 0.120000			Min Classification Error = 0.394000				
Probability of Error = 0.118500			Probability of Error = 0.101400				
Confusion Matrix =			Confusion Matrix =				
5581	206	66	75	5870	16	1	41
320	453	0	117	498	272	0	120
3	0	825	26	3	0	851	0
318	48	6	1956	314	20	1	1993
MLP MODEL 3 (500 training samples)			MLP MODEL 6 (5000 training samples)				
-----			-----				
Optimum number of perceptrons = 7			Optimum number of perceptrons = 10				
Min Classification Error = 0.146000			Min Classification Error = 0.400600				
Probability of Error = 0.113000			Probability of Error = 0.127400				
Confusion Matrix =			Confusion Matrix =				
5368	443	38	79	5871	25	3	29
98	729	0	63	566	23	0	301
0	0	854	0	0	0	854	0
293	113	3	1919	321	28	1	1978

In the graph shown below we plot the probability error obtained from each MLP compared to the theoretically optimal classifier. We see that the models have higher error at lower training set sizes but higher training set sizes could also be overfit to throw off the model.

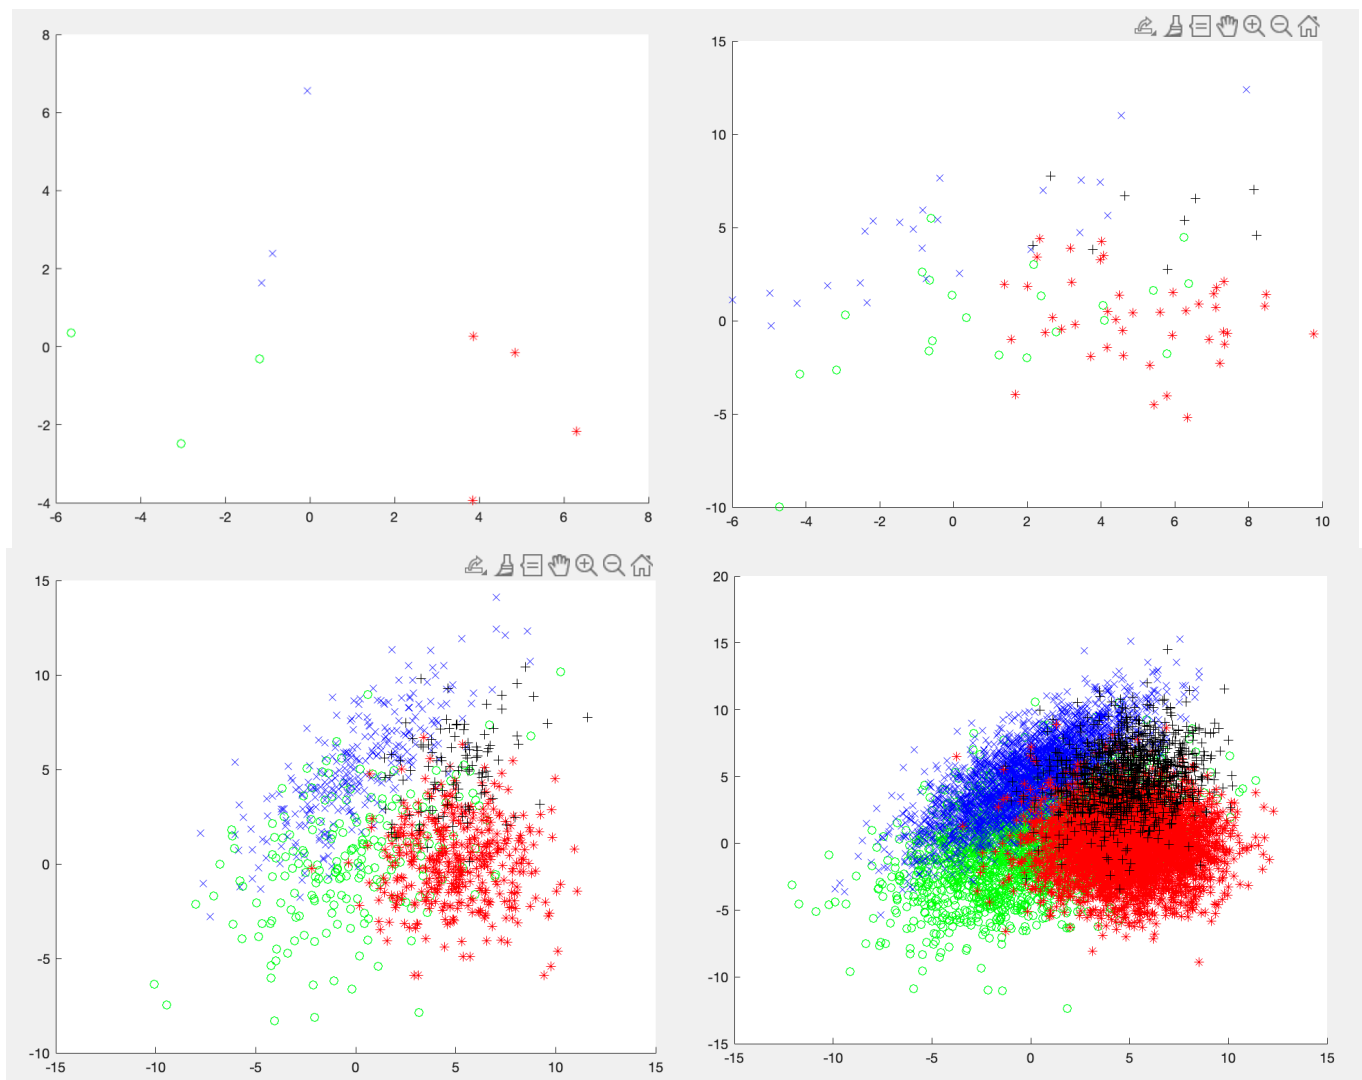


QUESTION 2

PROCESS

DATA GENERATION

Datasets comprising of 10, 100, 1000, and 10000 samples were created using a Gaussian Mixture Model that generates 2-dimensional data with four distinct classes, each having different class priors. The visual representations of the data distributions for these datasets are shown below:



The means and covariance matrices are given below:

$M1 = [0,0]$; $M2 = [0,5]$; $M3 = [5,0]$; $M4 = [5,5]$;
 $S1 = [10,4;4,10]$; $S2 = [8,6;6,8]$; $S3 = [5,0;0,5]$; $S4 = [4,1;1,5]$;

GAUSSIAN MIXTURE MODEL

The Gaussian Mixture Model is initialized with a random mean, variance and prior probabilities and using Gaussian Components ranging from 1 to 6 we fit the GMM distribution onto the training data and then calculate the log likelihood of the validation data. The max iterations was set at 3000, the tolerance set to 1×10^{-6} and the regularization value was given as 0.01.

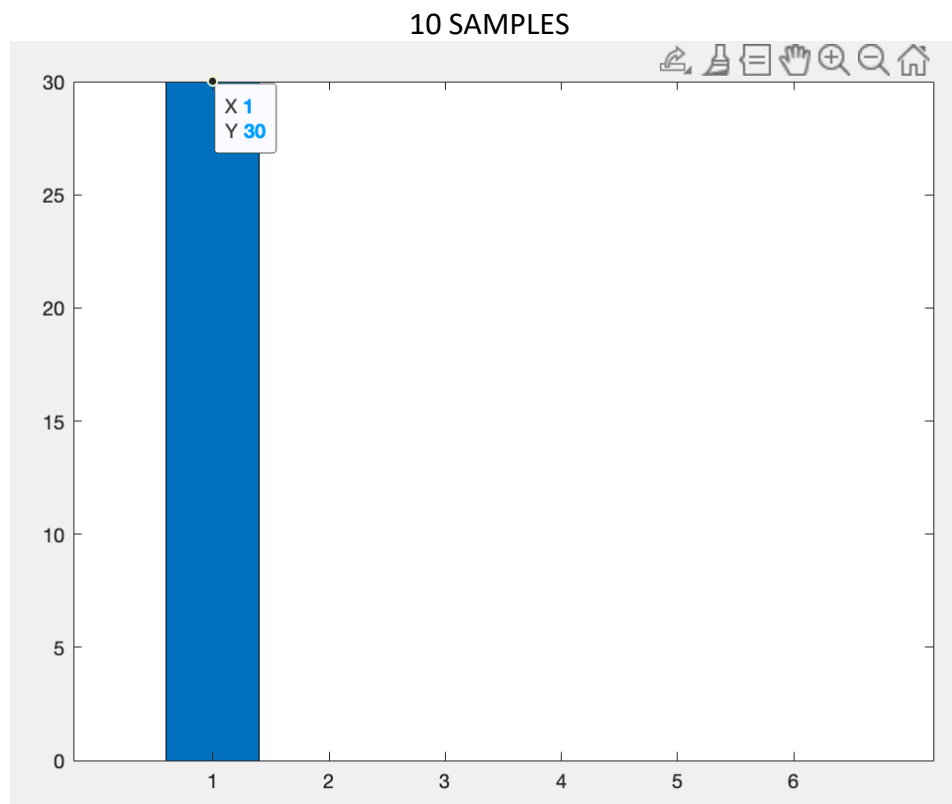
TEN-FOLD CROSS VALIDATION

Each GMM order with Gaussian components ranging from 1 to 6 is cross validated by splitting the data into 10 folds, taking each fold as a validation fold and the remaining as the training folds, we fit the GMM distribution onto the training data and then calculate the log likelihood of the validation data. The likelihood across all 10 folds is averaged to give the score for that model.

MODEL ORDER SELECTION


The model with the best score computed above will be selected as the best GMM model. This is repeated 30 times for each data set and the selection rates for each model are tallied.

RESULTS




MODEL SELECTED = 1

100 SAMPLES




MODEL SELECTED = 2

1000 SAMPLES



MODEL SELECTED = 4

10000 SAMPLES



MODEL SELECTED = 5

OBSERVATIONS

The Gaussian Mixture Model struggles to accurately classify data when there is overlap between clusters, especially with smaller sample sizes, and the convergence of the model can be time-consuming, requiring too many iterations which may result in incorrect classifications. Additionally, the hyperparameters of the model must be carefully tuned to produce effective results, otherwise the model could be susceptible to over-fitting.

APPENDIX

CODE FOR QUESTION 1 (MATLAB)

```
clc; clear all; close all;
```

```
n = 3;  
CL = 4;
```

```
priori = [0.25,0.25,0.25,0.25];  
mean_ij(1,:) = [0,0,0];  
mean_ij(2,:) = [0,0,5];  
mean_ij(3,:) = [0,5,0];  
mean_ij(4,:) = [5,0,0];  
cov_m(:, :, 1) = [1,0,-2;0,1,0;-2,0,12];  
cov_m(:, :, 2) = [6,0,0;0,0.3,0;0,0,0.3];  
cov_m(:, :, 3) = [1,0,-2;0,0.3,0;-2,0,12];  
cov_m(:, :, 4) = [6,0,2;0,1,0;2,0,12];
```

```
DS_train = [100,200,500,1000,2000,5000];  
DS_val = 10000;
```

```
%Perceptron Range  
PVal = 1:10;
```

```
%Learning Rate and Epochs for finding best number of Perceptrons  
alpha1 = 0.01;  
epochs1 = 100;
```

```
%Learning Rate and Epochs for training MLP  
alpha2 = 0.01;  
epochs2 = 100;
```

```
%Generate Test Data  
DSval = generate_samples(mean_ij,cov_m,priori,DS_val,n);  
Xval = DSval(:,1:n);  
mu_n = mean(Xval);  
sd_n = std(Xval);  
Xval = (Xval - repmat(mu_n, DS_val, 1)) ./ repmat(sd_n, DS_val, 1);  
Yval = DSval(:,CL) + 1;
```

```
for t = 1:numel(DS_train)  
    %Find Optimum Number of Perceptrons  
    %Generate Training Data  
    DS_size = DS_train(t);  
    Dtrain = generate_samples(mean_ij,cov_m,priori,DS_size,n);  
    fprintf('MLP MODEL %d (%d training samples)\n',t,DS_size);  
    fprintf('-----\n');  
    % Shuffle the rows of the data randomly  
    shf_idx = randperm(size(Dtrain, 1));  
    Dtrain_shf = Dtrain(shf_idx, :);  
  
    % Separate the input data from the class labels  
    X = Dtrain_shf(:,1:end-1);  
  
    mu_n = mean(X);  
    sd_n = std(X);  
    X = (X - repmat(mu_n, DS_size, 1)) ./ repmat(sd_n, DS_size, 1);  
    Y = Dtrain_shf(:,end);
```

```

Y = Y + 1;

% Convert the class labels to binary vectors
CL = max(Y);
T = zeros(length(Y),CL);
for i = 1:length(Y)
    T(i,Y(i)) = 1;
end

%Splitting the data into 10 folds
k = 10;
f_size = floor(DS_size/k);
Xfolds = zeros(f_size,n,k);
for fold = 1:k
    fold_start = (fold-1)*f_size + 1;
    fold_end = fold*f_size;
    Xfolds(:, :, fold) = X(fold_start:fold_end, :);
end

Tfolds = zeros(f_size,n+1,k);
for fold = 1:k
    fold_start = (fold-1)*f_size + 1;
    fold_end = fold*f_size;
    Tfolds(:, :, fold) = T(fold_start:fold_end, :);
end

% Set up the neural network with ELU activation function and SoftMax output layer
min_ce = 1;
for i = 1:length(PVal)
    inp = size(X,2);
    numH = PVal(i);
    outp = CL;

    % Evaluate the performance of the neural network using 10-fold cross-validation
    fold_ce = zeros(k,1);
    for fold = 1:k
        % Divide the data into training and validation sets
        Xval_fold = squeeze(Xfolds(:, :, fold));
        Tval_fold = squeeze(Tfolds(:, :, fold));

        t_idx = [1:(fold-1)*f_size, fold*f_size+1:DS_size];
        Xtrain_fold = X(t_idx, :);
        Ttrain_fold = T(t_idx, :);
        r_ce = zeros(10,1);
        best_ll = -Inf;
        for j = 1:10
            % Train the neural network on the training set
            wt1 = randn(inp,numH);
            bs1 = randn(numH,1);
            wt2 = randn(numH,outp);
            bs2 = randn(outp,1);
            for epoch = 1:epochs1
                % Forward propagation
                Z1 = Xtrain_fold*wt1 + repmat(bs1',size(Xtrain_fold,1),1);
                A1 = elu(Z1,alpha1);
                Z2 = A1*wt2 + repmat(bs2',size(Xtrain_fold,1),1);
                Y = s_max(Z2);

                % Backward propagation
                dZ2 = Y - Ttrain_fold;
                dwt2 = A1' * dZ2;
            end
        end
    end
end

```

```

        dbs2 = sum(dZ2,1)';
        dA1 = dZ2 * wt2';
        dZ1 = eluG(Z1) .* dA1;
        dwt1 = Xtrain_fold' * dZ1;
        dbs1 = sum(dZ1,1)';

        % Update the weights and biases of the neural network
        wt1 = wt1 - alpha1*dwt1;
        bs1 = bs1 - alpha1*dbs1;
        wt2 = wt2 - alpha1*dwt2;
        bs2 = bs2 - alpha1*dbs2;
        ytrain_fold = zeros(size(Tval_fold,1),1);
        for s = 1:size(Tval_fold,1)
            for c = 1:CL
                if Ttrain_fold(s,c) == 1
                    ytrain_fold(s) = c;
                end
            end
        end
        train_ll(epoch) = sum(log(Y(ytrain_fold',1)));
    end
    if (train_ll(end) > best_ll)
        best_ll = train_ll(end);
        rwt1 = wt1;
        rwt2 = wt2;
        rbs1 = bs1;
        rbs2 = bs2;
    end
end
% Evaluate the performance of the trained neural network on the validation
set
Z1 = Xval_fold*rwt1 + repmat(rbs1',size(Xval_fold,1),1);
A1 = elu(Z1,alpha1);
Z2 = A1*rwt2 + repmat(rbs2',size(Xval_fold,1),1);
Y = s_max(Z2);
[~,pred] = max(Y,[],2);
corr = 0;
for a = 1: numel(pred)
    p = pred(a);
    if (Tval_fold(a,p) == 1)
        corr = corr + 1;
    end
end
fold_ce(fold) = 1 - (corr/numel(pred));
end

% Calculate the average error over all folds
avg_fold_ce = mean(fold_ce);

% Update the minimum error and the corresponding neural network parameters
if avg_fold_ce <= min_ce
    min_ce = avg_fold_ce;
    bestwt1 = rwt1;
    bestbs1 = rbs1;
    bestwt2 = rwt2;
    bestbs2 = rbs2;
    bestnumH = numH;
end
end
best_P(t) = bestnumH;
fprintf('Optimum number of perceptrons = %d\n',bestnumH);

```

```

fprintf('Min Classification Error = %f\n',min_ce);

%Train MLP models using optimum number of perceptrons for each set
Xtrain = Dtrain(:,1:n);
mu_n = mean(Xtrain);
sd_n = std(Xtrain);
Xtrain = (Xtrain - repmat(mu_n, DS_size, 1)) ./ repmat(sd_n, DS_size, 1);
Ytrain = Dtrain(:,n+1) + 1;
Ttrain = zeros(length(Ytrain),CL);
for i = 1:length(Ytrain)
    Ttrain(i,Ytrain(i)) = 1;
end
inp = size(Xtrain,2);
outp = CL;
numH = best_P(t);
best_ll = -Inf;
for j = 1:10
    wt1 = randn(inp,numH);
    bs1 = zeros(numH,1);
    wt2 = randn(numH,outp);
    bs2 = zeros(outp,1);
    train_ll = zeros(epochs2, 1);
    for epoch = 1:epochs2
        % Forward pass
        Z1 = Xtrain*wt1 + repmat(bs1',DS_size,1);
        A1 = sig(Z1);
        Z2 = A1*wt2 + repmat(bs2',DS_size,1);
        Y = s_max(Z2);

        % Compute cross-entropy loss
        loss = -sum(log(Y(Ytrain',1)));

        % Backward pass
        dZ2 = Y - Ttrain;
        dwt2 = A1' * dZ2;
        dbs2 = sum(dZ2,1)';
        dA1 = dZ2 * wt2';
        dZ1 = sigG(Z1) .* dA1;
        dwt1 = Xtrain' * dZ1;
        dbs1 = sum(dZ1,1)';
        wt1 = wt1 - alpha2 * dwt1;
        bs1 = bs1 - alpha2 * dbs1;
        wt2 = wt2 - alpha2 * dwt2;
        bs2 = bs2 - alpha2 * dbs2;
        train_ll(epoch) = sum(log(Y(Ytrain',1)));
    end
    if (train_ll(end) > best_ll)
        best_ll = train_ll(end);
        bestwt1 = wt1;
        bestwt2 = wt2;
        bestbs1 = bs1;
        bestbs2 = bs2;
    end
end

%Test each MLP with the test data
Z1 = Xval*bestwt1 + repmat(bestbs1',DS_val,1);
A1 = sig(Z1);
Z2 = A1*bestwt2 + repmat(bestbs2',DS_val,1);
Y = s_max(Z2);

```

```

%Calculate Error and Confusion matrix
[~,pred] = max(Y,[],2);
conf_m = zeros(CL);
for i = 1:numel(pred)
    a = Yval(i);
    p = pred(i);
    conf_m(a,p) = conf_m(a,p) + 1;
end
error = 1 - (trace(conf_m)/DS_val);
fprintf('Probability of Error = %f\n',error);
fprintf('Confusion Matrix = \n');
disp(conf_m);
errors(t) = error;
end

%Theoretical Classifier
%Generate Test Data
DSval = generate_samples(mean_ij,cov_m,priori,DS_val,n);
Xval = DSval(:,1:n);
mu_n = mean(Xval);
sd_n = std(Xval);
Xval = (Xval - repmat(mu_n, DS_val, 1)) ./ repmat(sd_n, DS_val, 1);
Yval = DSval(:,CL) + 1;

%Loss matrix
loss_m = ones(CL) - eye(CL);

%Class-Conditional Probabilities
probX_L = zeros(CL,DS_val);
for i = 1:CL
    probX_L(i,:) = mvnpdf(Xval,mean_ij(i,:),squeeze(cov_m(:,:,i)));
end

%Class Posteriors
probX = priori * probX_L;
Class_Pos = (probX_L .* repmat(priori', 1, DS_val)) ./ repmat(probX, CL, 1);
exp_risk = loss_m * Class_Pos;
 [~,dec] = min(exp_risk,[],1);
dec = dec';

%Calculate Risk and Confusion Matrix
avg_exp_risk = sum(min(exp_risk,[],1))/DS_val;
opt_conf_m = zeros(CL);
for i = 1:CL
    for j = 1:CL
        opt_conf_m(i,j) = numel(find((i == Yval) & (j == dec)));
    end
end

% Plot graph to compare theoretical classifier and various MLP models
plot(DS_train, errors, '-', 'MarkerSize', 15);
hold on;
yline(avg_exp_risk, 'r-', 'LineWidth', 1);
xlabel('Number of Samples');
ylabel('Error');
title('Plot to Compare Errors of the trained models');
hold off;

function DS = generate_samples(mean_ij,cov_m,priori,DS_size,n)
    Priori_Cum = cumsum(priori);
    rand = randn(DS_size, 1);

```

```

CL = zeros(size(rand));
for i = 1:DS_size
    if rand(i) <= Priori_Cum(1)
        CL(i) = 0;
    elseif rand(i) <= Priori_Cum(2)
        CL(i) = 1;
    elseif rand(i) <= Priori_Cum(3)
        CL(i) = 2;
    else
        CL(i) = 3;
    end
end
DS = zeros(DS_size,n);
for i = 1:DS_size
    if CL(i) == 0
        DS(i,:) = mvnrnd(mean_ij(1,:),squeeze(cov_m(:,:,1)));
    elseif CL(i) == 1
        DS(i,:) = mvnrnd(mean_ij(2,:),squeeze(cov_m(:,:,2)));
    elseif CL(i) == 2
        DS(i,:) = mvnrnd(mean_ij(3,:),squeeze(cov_m(:,:,3)));
    elseif CL(i) == 3
        DS(i,:) = mvnrnd(mean_ij(4,:),squeeze(cov_m(:,:,4)));
    end
end
DS = [DS,CL];
end

function y = s_max(x)
    m = max(x,[],2);
    y = exp(x - m) ./ sum(exp(x - m),2);
end

function y = elu(x,alpha)
    y = x .* (x > 0) + alpha * (exp(x) - 1) .* (x <= 0);
end

function y = eluG(x)
    y = ones(size(x));
    y(x < 0) = exp(x(x < 0));
end

function y = sig(x)
    % Sigmoid activation function
    y = 1 ./ (1 + exp(-x));
end

function y = sigG(x)
    % Computes the gradient of the sigmoid function at x
    y = sig(x).*(1-sig(x));
end

```

CODE FOR QUESTION 2 (MATLAB)

```
clc; clear all; close all;
```

```
n_features = 2;  
n_comp = 4;  
priori = [0.2,0.3,0.4,0.1];  
comp_mu(1,:) = [0,0];  
comp_mu(2,:) = [0,5];  
comp_mu(3,:) = [5,0];  
comp_mu(4,:) = [5,5];  
comp_sd(:,:,1) = [10,4;4,10];  
comp_sd(:,:,2) = [8,6;6,8];  
comp_sd(:,:,3) = [5,0;0,5];  
comp_sd(:,:,4) = [4,1;1,5];  
set_size = [10,100,1000,10000];  
folds = 10;  
g_comp = 6;
```

```
select = zeros(numel(set_size),g_comp);
```

```
for s = 1:numel(set_size)
```

```
    %Generate Data
```

```
    n_samples = set_size(s);
```

```
    DS = generate_samples(comp_mu,comp_sd,priori,n_samples,n_features,n_comp);
```

```
    plot_data(DS);
```

```
    n_repeat = 30;
```

```
for rep = 1:n_repeat
```

```
    %Shuffle Data
```

```
    shf_idx = randperm(n_samples);
```

```
    DS = DS(shf_idx, :);
```

```
    %Split data into X and Y
```

```
    X = DS(:,1:n_features);
```

```
    Y = DS(:,n_features+1);
```

```
    %Split data into 10 folds
```

```
    n_folds = 10;
```

```
    f_size = floor(n_samples/n_folds);
```

```
    Xfolds = zeros(f_size,n_features,n_folds);
```

```
    for f = 1:n_folds
```

```
        fold_start = (f-1)*f_size + 1;
```

```
        fold_end = f*f_size;
```

```
        Xfolds(:,:,f) = X(fold_start:fold_end,:);
```

```
    end
```

```
    %Perform 10-fold cross validation for each number of components
```

```
    max_ll = 0;
```

```
    comp_ll = zeros(g_comp,1);
```

```
    for g = 1:g_comp
```

```
        fold_ll = zeros(n_folds,1);
```

```
        for f = 1:n_folds
```

```
            %Assign Training and Validation Folds
```

```
            Xval_fold = squeeze(Xfolds(:,:,f));
```

```
            t_idx = [1:(f-1)*f_size,f*f_size+1:n_samples];
```

```
            Xtrain_fold = X(t_idx,:);
```

```
            % Estimate GMM parameters using EM algorithm
```

```

        gmFit = fitgmdist(Xtrain_fold, g, 'RegularizationValue', 0.01,
'ProbabilityTolerance',1e-6, 'Options', statset('MaxIter', 1000));

        % Evaluate log-likelihood of validation set
        fold_ll(f) = sum(log(pdf(gmFit,Xval_fold)));
    end
    comp_ll(g) = mean(fold_ll);
end
[max_ll,g_select] = max(comp_ll);
select(s,g_select) = select(s,g_select) + 1;
end
fprintf('Set with %d samples completed.\n',n_samples);
figure
bar(select(s,:))
end

function DS = generate_samples(comp_mu,comp_sd,priori,n_samples,n,c)

x = zeros(n, n_samples);
labels = zeros(1, n_samples);

u = rand(1, n_samples);
th = zeros(1, c+1);
th(1:c) = cumsum(priori);
th(c+1) = 1;

for l = 1:c
    indl = find(u <= th(l));
    Nl = length(indl);
    labels(indl) = (l-1)*ones(1, Nl);
    u(indl) = 1.1;
    x(:, indl) = mvnrnd(comp_mu(l, :), squeeze(comp_sd(:,:,l)), Nl)';
end
x = x';
labels = labels' + 1;
DS = [x,labels];
end

function plot_data(DS)
figure
scatter(DS(find(DS(:,3)==1),1),DS(find(DS(:,3)==1),2),'o', 'g')
hold on
scatter(DS(find(DS(:,3)==2),1),DS(find(DS(:,3)==2),2),'X', 'b')
scatter(DS(find(DS(:,3)==3),1),DS(find(DS(:,3)==3),2),'*', 'r')
scatter(DS(find(DS(:,3)==4),1),DS(find(DS(:,3)==4),2),'+', 'k')
end

```